

# Signals

---

# Signals

Signals are software interrupts.

Signals are a way for a process to be notified of asynchronous events.

Some examples:

- a timer you set has gone off (SIGALRM)
- some I/O you requested has occurred (SIGIO)
- a user resized the terminal  
"window" (SIGWINCH) a user  
disconnected from the system  
(SIGHUP)
- ...

See also: `signal(2)`/`signal(3)`/`signal(7)` (note: these man pages vary significantly across platforms!)

# Signals

Besides the asynchronous events listed previously, there are many ways to generate a signal:

- terminal generated signals (user presses a key combination which causes the terminal driver to generate a signal)
- hardware exceptions (divide by 0, invalid memory references, etc)
- kill(1) allows a user to send any signal to any process (if the user is the owner or superuser)
- kill(2) (a system call, not the unix command) performs the same task
- software conditions (other side of a pipe no longer exists, urgent data has arrived on a network file descriptor, etc.)

# Signal Concepts

Once we get a signal, we can do one of several things:

1. Ignore it. (note: there are some signals which we CANNOT or SHOULD NOT ignore)
2. Catch it. That is, have the kernel call a function which we define whenever the signal occurs.
3. Accept the default. Have the kernel do whatever is defined as the default action for this signal

# Signal Concepts

Name	Description	Default action
<code>SIGABRT</code>	abnormal termination ( <code>abort</code> )	terminate+core
<code>SIGALRM</code>	timer expired ( <code>alarm</code> )	terminate

# signal(3)

---

```
#include <signal.h>
```

```
void (*signal(int signo, void (*func)(int)))(int);
```

Returns: previous disposition of signal if OK, SIG\_ERR otherwise

-

*func* can be:

- SIG\_IGN which requests that we ignore the signal *signo*
  - SIG\_DFL which requests that we accept the default action for signal *signo*
  - or the address of a function which should catch or handle a signal (signal handler or the signal-catching function)
-

# signal(3)

---

- The prototype for the signal function states that the function requires two arguments and returns a pointer to a function that returns nothing (void).
  - The signal function's first argument, *signo*, is an integer.
  - The second argument is a pointer to a function that takes a single integer argument and returns nothing.
  - The function whose address is returned as the value of signal takes a single integer argument (the final (int)).
  - i.e., this declaration says that the signal handler is passed a
  - single integer argument (the signal number) and that it returns nothing. When we call signal to establish the signal handler, the second argument is a pointer to the function.
  - The return value from signal is the pointer to the previous signal handler.
-

# sigaction(2)

---

```
#include <signal.h>
```

```
int sigaction(int signo, const struct sigaction *act,  
struct sigaction *oldact);
```

This function allows us to examine or modify the action associated with a particular signal.

```
struct sigaction {  
    void (*sa_handler)();           /* addr of signal handler, or  
                                    SIG_IGN or SIG_DFL */  
    sigset_t sa_mask;               /* additional signals to block */  
    int sa_flags;                   /* signal options */  
};
```

signal(3) is (nowadays) commonly implemented via sigaction(2).

---



# sigprocmask function

the signal mask of a process is the set of signals currently blocked from delivery to that process.

A process can examine its signal mask, change its signal mask, or perform both operations in one step by calling the following function.

```
#include <signal.h>

int sigprocmask(int how, const sigset_t *restrict set,
                sigset_t *restrict oset);
```

Returns: 0 if OK, 1 on error

# sigprocmask function

## Ways to change current signal mask using `sigprocmask`

<i>how</i>	Description
<code>SIG_BLOCK</code>	The new signal mask for the process is the union of its current signal mask and the signal set pointed to by <i>set</i> . That is, <i>set</i> contains the additional signals that we want to block.
<code>SIG_UNBLOCK</code>	The new signal mask for the process is the intersection of its current signal mask and the complement of the signal set pointed to by <i>set</i> . That is, <i>set</i> contains the signals that we want to unblock.
<code>SIG_SETMASK</code>	The new signal mask for the process is replaced by the value of the signal set pointed to by <i>set</i> .

## kill(2) and raise(3)

---

```
#include <sys/types.h> #include  
<signal.h>  
  
int kill(pid_t pid, int signo); int raise(int  
signo);
```

- $pid > 0$  – signal is sent to the process whose PID is  $pid$
  - $pid == 0$  – signal is sent to all processes whose process group ID equals the process group ID of the sender
  - $pid == -1$  – POSIX.1 leaves this undefined, BSD defines it (see kill(2))
-

## More advanced signal handling via signal sets

---

- `int sigemptyset(sigset_t *set)` – initialize a signal set to be empty
- `int sigfillset(sigset_t *set)` – initialize a signal set to contain all signals
- `int sigaddset(sigset_t *set, int signo)` `int`
- `sigdelset(sigset_t *set, int signo)` `int`
- `sigismember(sigset_t *set, int signo)`

–

## Interrupted System Calls

---

Some system calls can block for long periods of time (or forever). These include things like:

- `read(2)`s from files that can block (pipes, networks, terminals)
- `write(2)` to the same sort of files
- `open(2)` of a device that waits until a condition occurs (for example, a modem)
- `pause(3)`, which purposefully puts a process to sleep until a signal occurs
- certain `ioctl(3)`s
- certain IPC functions

Catching a signal during execution of one of these calls traditionally led to the process being aborted with an `errno` return of `EINTR`.

---

## Interrupted System Calls

Previously necessary code to handle

---

EINTR: again:

```
    if ((n = read(fd, buf, BUFSIZE)) < 0) { if
        (errno == EINTR)
            goto again; /* just an interrupted system call */
        /* handle other errors */
    }
```

Nowadays, many Unix implementations automatically restart certain system calls.

---