

# Programming with UNIX File Systems

(Chap 3, 4. in the book “Advanced Programming in the UNIX Environment”)

# File

- A file is a contiguous sequence of bytes.
- No format imposed by the operating system.
- Each byte is individually addressable in a disk file.

# File Descriptor

- **open ( )** returns an **fd**, an integer value.
- Used in subsequent I/O operations on that file.
- **close (fd)** closes that file described by **fd**.
- All of a process's open files are automatically closed when it terminates.

# File Descriptor

- file descriptor: 0 ~ N (N = 19? or more) // `unistd.h`

Value	Meaning
0	standard input
1	standard output
2	standard error
3 .. 19	fds for users

# System call for file I/O

- `open()`
- `creat()`
- `close()`
- `read()`
- `write()`
- `lseek()`

# open()

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int open (const char *pathname, int oflag, [ mode_t mode ]);
```

- Open a file
  - The **open** API is used to create new files and also to establish a connection between a process and a file
- A new file can be created if there is no such file
  - pathname : file name
  - mode : access permission. Can be omitted
  - return value : file descriptor value. if fail, return -1

# open()

- Second parameter `o_flag` is (logical) OR of following constants

- Ex) 

```
int fd;  
fd = open("afile", O_RDWR | O_CREAT, 0644);
```

- At least one of following constants need to be specified

- `O_RDONLY` : for read only. No write can be performed.
- `O_WRONLY` : for write only. No read can be performed.
- `O_RDWR` : for read or write, or both

# open()

- Following constants for o\_flag is optional
  - O\_APPEND : append to the end of file on each write.
  - O\_CREAT : create the file if it doesn't exist. (mode is applied)
  - O\_EXCL : generate an error if O\_CREAT is also specified and the file already exists.
  - O\_TRUNC : if the file exists, and if the file is successfully opened for either write-only or read-write, truncate its length to 0.
  - O\_SYNC : Have each write wait for physical I/O to complete



# umask ( ) :

- The umask function set up default permissions for newly created files.
- Syntax:

```
#include <sys/stat.h>  
mode_t umask(mode_t cmask);
```

The file mode creation mask is used whenever the process creates a new file or a new directory.

**chown** command is used to change the file permissions. These permissions read, write and execute permission for owner, group, and others.

### Syntax

```
#include<unistd.h>
```

```
int chown(const char * pathname , uid_t owner , gid_t group);
```

- The first optional parameter indicates who – this can be (u)ser, (g)roup, (o)thers or (a)ll.
- The second optional parameter indicates opcode – this can be for adding (+), removing (-) or assigning (=) a permission.
- The third optional parameter indicates the mode – this can be (r)ead, (w)rite, or e(x)ecute.

**Example:** chown owner-user file

```
chown owner-user:owner-group file
# ls -l demo.txt
-rw-r--r-- 1 root root 0 Aug 31 05:48 demo.txt

# chown vivek demo.txt

# ls -l demo.txt
-rw-r--r-- 1 vivek root 0 Aug 31 05:48 demo.txt
```

# creat()

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int creat ( const char *pathname, mode_t mode );
```

- Create a new file
  - pathname : file name
  - mode : access permission
  - Return value : file descriptor. If failed, return -1

# creat()

## ■ Following two lines are functionally the same

- `fd = creat ( pathname, mode );`
- `fd = open ( pathname, O_WRONLY | O_CREAT | O_TRUNC, mode);`

## ■ mode

- 0644 ☐ -rw-r--r--
- 0755 ☐ -rwxr-xr-x
- 0444 ☐ -r--r--r--

# close()

```
#include <unistd.h>

int close ( int filedesc );
```

- Close a file
- If a process is terminated, all open files will be automatically closed.

# read()

```
#include <unistd.h>
ssize_t read ( int filedes, void *buf, size_t nbytes );
```

- Read data from a file
- **buf** : memory address to store the data to be read
- **nbytes** : the number of data bytes to be read
- Return value
  - On success, the number of data bytes read
  - If the end of file, 0
  - If failed, -1
- **size\_t** : unsigned int

# Example : count.c

```
#include <stdlib.h> // count.c
#include <fcntl.h>
#include <unistd.h>
#define BUFSIZE 512

int main()
{
    char buffer[BUFSIZE];
    int filedес;
    ssize_t nread;
    long total = 0;

    if ((filedes = open("afile", O_RDONLY)) == -1)
    { perror("afile"); exit(-1); }

    while( (nread = read(filedes, buffer, BUFSIZE)) > 0)
        total += nread;

    close(filedes);
    printf ("total chars in afile: %ld\n", total);
    return 0;
}
```

# write()

```
#include <unistd.h>
ssize_t write (int filedes, const void *buf, size_t nbytes);
```

- Write data to a file
- **buf** : memory address to write data
- **nbytes** : number of data bytes to write
- Return value
  - On success : number of bytes written
  - -1 if failed.



# lseek()

```
#include <sys/types.h>
#include <unistd.h>
off_t lseek (int  filedes,  off_t  offset, int  whence );
```

- Move current file offset
- Parameters
  - **Whence** : start point to move from
    - SEEK\_SET
    - SEEK\_CUR
    - SEEK\_END
  - **Offset** : the number of bytes to move from whence (negative value is possible)
  - Return value : new file offset if OK, -1 on error

# lseek()

- Current file offset
  - File read and write are performed at the point of current file offset
  - When a file is opened, current file offset is 0
  - After file read/write, file offset is automatically updated(moved).
  - We can move current file offset to an arbitrary location.

# Example : lseek1.c

```
#include <unistd.h> /* lseek1.c */
#include <fcntl.h>

char buf1[] = "abcdefghij";
char buf2[] = "ABCDEFGH IJ";

int main() {
    int fd;

    if ( (fd = creat("file.hole", 0644)) < 0)
    { perror("file.hole"); exit(-1); }

    if (write(fd, buf1, 10) != 10) /* offset now = 10 */
        perror("buf1");

    if (lseek(fd, 40, SEEK_SET) == -1) /* offset now = 40 */
        perror("lseek");

    if (write(fd, buf2, 10) != 10) /* offset now = 50 */
        perror("buf2");

    return 0;
}
```

# lseek1.c

## ■ Output “file.hole”

	0	1	2	3	4	5	6	7	8	9
0	a	b	c	d	e	f	g	h	i	j
10	\0	\0	\0	\0	\0	\0	\0	\0	\0	\0
20	\0	\0	\0	\0	\0	\0	\0	\0	\0	\0
30	\0	\0	\0	\0	\0	\0	\0	\0	\0	\0
40	A	B	C	D	E	F	G	H	I	J

# Files & Directories

## stat() – returns information about a file

```
#include <sys/types.h>
#include <sys/stat.h>
int  stat (const char *pathname, struct stat *buf );
int  fstat (int filedes, struct stat *buf );
int  lstat (const char *pathname, struct stat *buf );
```

- put information about a file into a stat type variable “buf”
- lstat() returns information about the symbolic link, not the file referenced by the symbolic link if the give file is a symbolic link.
- Return 0 on success. -1 if failed.

## stat structure (1/2)

- Defined in <sys/stat.h>

```
struct stat {  
    mode_t  st_mode;    /* file type & mode (permissions) */  
    ino_t    st_ino;     /* i-node number (serial number) */  
    dev_t    st_dev;     /* device number (filesystem) */  
    dev_t    st_rdev;    /* device number for special files */  
    nlink_t  st_nlink;   /* number of links */  
    uid_t    st_uid;     /* user ID of owner */  
    gid_t    st_gid;     /* group ID of owner */  
    off_t    st_size;    /* size in bytes, for regular files */  
    time_t   st_atime;    /* time of last access */  
    time_t   st_mtime;    /* time of last modification */  
    time_t   st_ctime;    /* time of last file status change */  
    long     st_blksize; /* best I/O block size */  
    long     st_blocks;  /* number of 512-byte blocks allocated */  
};
```

# File Types (1/3)

- Regular File
  - The most common type of file
  - Contains data of some form.
- Directory File
  - A file that contains the names of other files and pointers to information on these files.
- Character Special File
  - A type of file used for certain types of devices on a system
  - A device transmitting data character by character (**c**-----)
- Block Special File
  - Typically used for disk devices.
  - A device transmitting data block by block (**b**-----)



## File Types (2/3)

- Special File


```
cs1.kangwon.ac.kr - Zterm
cs1 [ ysmoon {95} /devices/pci@1f,2000/scsi@1 ]
cs1 [ ysmoon {95} /devices/pci@1f,2000/scsi@1 ] pwd
/devices/pci@1f,2000/scsi@1
cs1 [ ysmoon {96} /devices/pci@1f,2000/scsi@1 ] ls -l sd@1,0:*
brw-r----- 1 root sys 32, 128 3 7 14:42 sd@1,0:a
crw-r----- 1 root sys 32, 128 3 7 14:42 sd@1,0:a,raw
brw-r----- 1 root sys 32, 129 3 7 14:42 sd@1,0:b
crw-r----- 1 root sys 32, 129 3 7 14:42 sd@1,0:b,raw
brw-r----- 1 root sys 32, 130 3 7 14:42 sd@1,0:c
crw-r----- 1 root sys 32, 130 3 7 14:42 sd@1,0:c,raw
brw-r----- 1 root sys 32, 131 3 7 14:42 sd@1,0:d
crw-r----- 1 root sys 32, 131 3 7 14:42 sd@1,0:d,raw
brw-r----- 1 root sys 32, 132 3 7 14:42 sd@1,0:e
crw-r----- 1 root sys 32, 132 3 7 14:42 sd@1,0:e,raw
brw-r----- 1 root sys 32, 133 3 7 14:42 sd@1,0:f
crw-r----- 1 root sys 32, 133 3 7 14:42 sd@1,0:f,raw
brw-r----- 1 root sys 32, 134 3 7 14:42 sd@1,0:g
crw-r----- 1 root sys 32, 134 3 7 14:42 sd@1,0:g,raw
brw-r----- 1 root sys 32, 135 2 27 21:38 sd@1,0:h
crw-r----- 1 root sys 32, 135 3 7 14:42 sd@1,0:h,raw
cs1 [ ysmoon {97} /devices/pci@1f,2000/scsi@1 ]
```

## File Types (3/3)

- **FIFO**
  - A type of file used for interprocess communication between processes
  - Sometimes called named pipe
- **socket**
  - A type of file used for network communication between processes
  - Typically used in Network Programming
- **Symbolic link**
  - A type of file that points to another file.

## Checking File Type (1/2)

- Macros for checking file type.
  - Defined in stat.h (/usr/include/sys/stat.h)
  - S\_ISREG() : regular file
  - S\_ISDIR() : directory file
  - S\_ISCHR() : character special file
  - S\_ISBLK() : block special file
  - S\_ISFIFO() : pipe or FIFO
  - S\_ISLNK() : symbolick link
  - S\_ISSOCK() : socket
- Return 1 if YES, otherwise return 0
- Testing st\_mode field in stat structure



## Checking File Type (2/2)

- File type constants

- Defined in stat.h (/usr/include/sys/stat.h)
- S\_IFREG : regular file
- S\_IFDIR : directory file
- S\_IFCHR : character special file
- S\_IFBLK : block special file
- S\_IFFIFO : pipe or FIFO
- S\_IFLNK : symbolic link
- S\_IFSOCK : socket

- S\_ISxxx() macro function checks whether st\_mode value has a bit of S\_IFxxx constant

- Ex) `#define S_ISDIR(mode) ((mode) & S_IFMT) == S_IFDIR)`

## example: stat.c (1/2)

```
#include <sys/types.h>    /* stat.c */
#include <sys/stat.h>
int  main(int argc, char *argv[])
{
    int  i;
    struct stat  buf;
    char  *ptr;
    for (i = 1; i < argc; i++) {
        printf("%s: ", argv[i]);
        if (lstat(argv[i], &buf) < 0) {
            perror("lstat()");  continue;
        }
        if      (S_ISREG(buf.st_mode))  ptr = "regular";
        else if (S_ISDIR(buf.st_mode))  ptr = "directory";
        else if (S_ISCHR(buf.st_mode))  ptr = "character special";
        else if (S_ISBLK(buf.st_mode))  ptr = "block special";
        else if (S_ISFIFO(buf.st_mode)) ptr = "fifo";
        else if (S_ISLNK(buf.st_mode))  ptr = "symbolic link";
        else if (S_ISSOCK(buf.st_mode)) ptr = "socket";
        else      ptr = "** unknown mode **";
        printf("%s\n", ptr);
    }
    exit(0);
}
```

## Example : stat.c (2/2)

```
$ a.out /etc /dev/ttya /bin a.out  
/etc: directory  
/dev/ttya: symbolic link  
/bin: symbolic link  
a.out: regular
```

- Use lstat() to get information of a symbolic link.

# File Permissions

- File access permission bits (st\_mode value in stat structure)

## st\_mode

type	special	permission
4 bits	3 bits	9 bits

st_mode mask	Meaning	Octal Code
S_IRUSR	user-read	0400
S_IWUSR	user-write	0200
S_IXUSR	user-execute	0100
S_IRGRP	group-read	0040
S_IWGRP	group-write	0020
S_IXGRP	group-execute	0010
S_IROTH	other-read	0004
S_IWOTH	other-write	0002
S_IXOTH	other-execute	0001

## Related UNIX Commands

- **chmod**

- Set File Access Permission
- Modify st\_mode value in stat structure

- **chown**

- Set User ID of file owner
- Modify st\_uid value in stat structure

- **chgrp**

- Set Group ID of file owner
- Modify st\_gid in stat structure



## Permissions (1/2)

- Read permission is necessary for
  - Opening a file with O\_RDONLY, O\_RDWR
- Write permission is necessary for
  - Opening a file with O\_WRONLY, O\_RDWR, O\_TRUNC
- write permission and execute permission in a directory is necessary for
  - Creating a file in the directory
  - Deleting a file in the directory

## Permissions (2/2)

- Whenever we want to open any type of file by name, we must have execute permission in each directory mentioned in the name, including the current directory if it is implied.
- For example, to open the file `/usr/dict/words`, we need execute permission in the directory `/`, `/usr`, `/usr/dict`, and then appropriate permission for the file `words`.
- In directory,
  - Read permission is necessary to read the list of files contained in the directory
  - Write permission is necessary for creating and deleting a file in the directory.
  - Execute permission is necessary for opening a file in the directory

# Effective User ID

- Real User ID/Real Group ID

- Who we really are
- Taken from our entry in the password file when we log in. normally these values don't change during a login session.

- Effective User ID/Effective Group ID

- Process's attribute
- Normally equals the real user ID /effective group id.
- However, can be different when S\_ISUID and S\_ISGID bit is set.

(setting the bit flag means “when this file is executed, set the effective user ID of the process to be the owner of the file (st\_uid))


- Can be understood as a user id/group id when a process is being executed.

## S\_ISUID , S\_ISGID (1/2)

- st\_mode bit in stat structure

- S\_ISUID : set-user-ID
- S\_ISGID : set-group-ID

### st\_mode



- When executing a file where st\_mode' S\_ISUID bit is set
  - effective user ID of the process is set to be the owner's user id of the file
- When executing a file where st\_mode's S\_ISGID bit is set
  - Effective group ID of the process is set to be the owner 's group id of the file

## S\_ISUID , S\_ISGID (2/2)

- When executing a file with S\_ISUID and S\_ISGID bits set
  - The file is executed with the permission of the file owner instead of the permission of Real User ID/Real Group ID
- example)

```
$ ls -al /bin/passwd
-r-sr-sr-x  1 root    sys      23500 Jul 30  2003 /bin/passwd*
```

## access()

```
#include <unistd.h>

int access (const char *pathname, int mode );
```

- Test the permission of a file with Real User ID and Real Group ID
- Return 0 if OK, -1 on error
- mode value :

mode	Description
<b>R_OK</b>	test for read permission
<b>W_OK</b>	test for write permission
<b>X_OK</b>	test for execute permission
<b>F_OK</b>	<i>test for existence of file</i>

## example: access()

```
#include <sys/types.h> // access.c
#include <fcntl.h>
#include <unistd.h>

int main(int argc, char *argv[])
{
    if (argc != 2) {
        printf("usage: a.out <pathname>\n");
        exit(-1);
    }

    if (access(argv[1], R_OK) < 0)
        perror("R_OK");
    else
        printf("read access OK\n");

    if (open(argv[1], O_RDONLY) < 0)
        perror("O_RDONLY");
    else
        printf("open for reading OK\n");

    return 0;
}
```

# chmod(), fchmod()

```
#include <sys/stat.h>
#include <sys/types.h>
int  chmod (const char *pathname, mode_t mode );
int  fchmod (int filedes, mode_t mode );
```

- Modify Access Permission of a File
  - Modify st\_mode value of stat structure
- Return 0 if OK, -1 on error
- mode : bitwise OR
  - S\_ISUID, S\_ISGID, S\_ISVTX
  - S\_IRUSR, S\_IWUSR, S\_IXUSR
  - S\_IRGRP, S\_IWGRP, S\_IXGRP
  - S\_IROTH, S\_IWOTH, S\_IXOTH



## example: chmod() (1/2)

```
#include <sys/types.h> // nchmod.c
#include <sys/stat.h>

int main()
{
    struct stat    statbuf;

    /* turn on both set-group-ID and group-execute */
    if (stat("foo", &statbuf) < 0)
        perror("stat(foo)");
    if (chmod("foo", (statbuf.st_mode | S_IXGRP | S_ISGID)) < 0)
        perror("chmod(foo)");

    /* set absolute mode to "rw-r--r--" */
    if (chmod("bar", S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH) < 0)
        perror("chmod(bar)");

    return 0;
}
```

## example: chmod() (2/2)

```
cs1.kangwon.ac.kr - Zterm 90x23
cs1 [ ysmoon {122} ~/unix/APUE ]
cs1 [ ysmoon {122} ~/unix/APUE ] touch bar foo
cs1 [ ysmoon {123} ~/unix/APUE ]
cs1 [ ysmoon {123} ~/unix/APUE ] chmod 0400 bar
cs1 [ ysmoon {124} ~/unix/APUE ] chmod 0744 foo
cs1 [ ysmoon {125} ~/unix/APUE ] ls -l bar foo
-r----- 1 ysmoon prof 0 3월 17일 20:24 bar
-rwxr--r-- 1 ysmoon prof 0 3월 17일 20:24 foo*
cs1 [ ysmoon {126} ~/unix/APUE ]
cs1 [ ysmoon {126} ~/unix/APUE ] gcc -o nchmod nchmod.c
cs1 [ ysmoon {127} ~/unix/APUE ]
cs1 [ ysmoon {127} ~/unix/APUE ] nchmod
cs1 [ ysmoon {128} ~/unix/APUE ]
cs1 [ ysmoon {128} ~/unix/APUE ] ls -l bar foo
-rw-r--r-- 1 ysmoon prof 0 3월 17일 20:24 bar
-rwxr--sr-- 1 ysmoon prof 0 3월 17일 20:24 foo*
cs1 [ ysmoon {129} ~/unix/APUE ]
```

# chown()

```
#include <sys/types.h>
#include <unistd.h>
int  chown (const char *pathname, uid_t owner, gid_t group );
int  fchown (int filedes, uid_t owner, gid_t group );
int  lchown (const char *pathname, uid_t owner, gid_t group );
```

- Modify User ID and Group ID of a file
  - Modify st\_uid, st\_gid value in stat structure
- Return 0 if OK, -1 on error
- lchown() modifies symbolic link itself
- Can be different in different versions of UNIX
  - On BSD-based system , only super-user can modify
  - On System V UNIX sysem, general user can modify.

## truncate(), ftruncate()

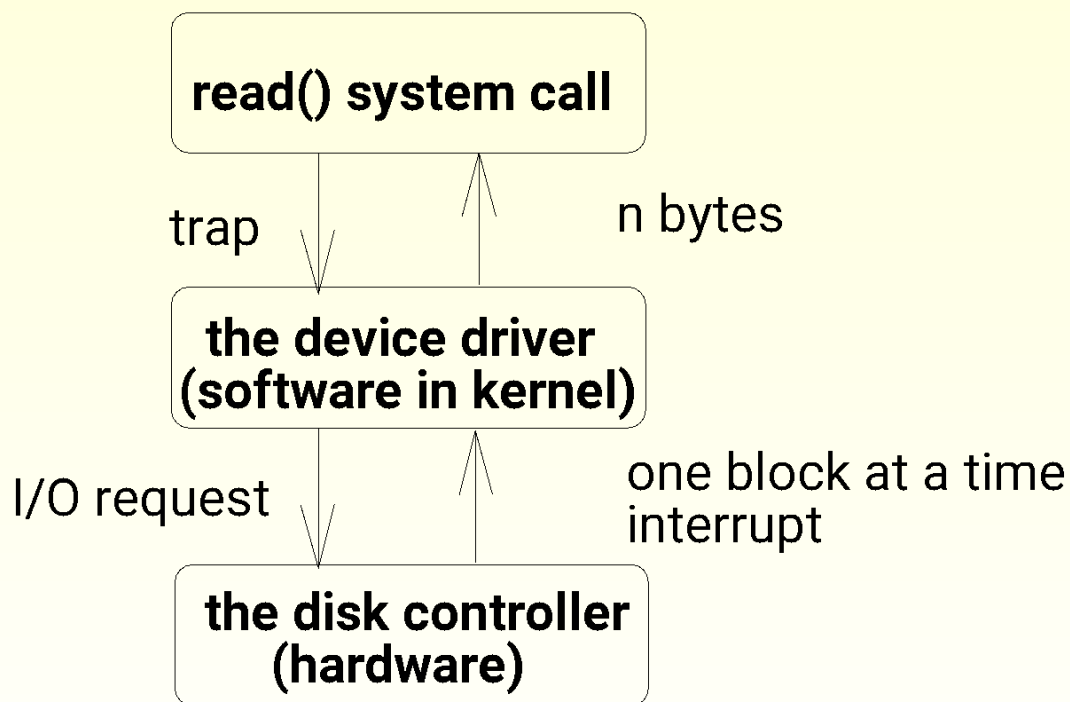
```
#include <sys/types.h>
#include <unistd.h>

int truncate (const char *pathname, off_t length );
int ftruncate (int filedes, off_t length );
```

- Reduce the size of a file into “length”
- Return 0 if OK, -1 on error

## Block I/O

- I/O is always done in terms of blocks.
- Sequence of a read() system call



## link()

```
#include <unistd.h>

int link (const char *existingpath, const char* newpath);
```

- Any file can have multiple directory entries pointing to its i-node.
- This function creates a link to an existing file
- Increment the link count

## unlink()

```
#include <unistd.h>

int  unlink (const char *pathname);
```

- Removes the directory entry and decrements the link count of the file.
- Return 0 if OK, -1 on error
- If the file is deleted, inode and data block become freed (deleted)

## example: unlink() (1/2)

```
#include <sys/types.h> // unlink.c
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int main() {
    int fd, len;
    char buf[20];
    fd = open("tempfile", O_RDWR | O_CREAT | O_TRUNC, 0666);
    if (fd == -1) perror("open1");
    close(fd);

    system("ls -l");

    unlink("tempfile");

    system("ls -l");

    return 0;
}
```



## example: unlink() (2/2)

```
cs1.kangwon.ac.kr - Zterm
cs1 [ ysmoon {152} ~/unix/APUE ] gcc -o unlink unlink.c
cs1 [ ysmoon {153} ~/unix/APUE ] unlink
ls -l
-rwxr-xr-x  1 ysmoon  prof      6452  3  17  20:24 nchmod
-rw-r--r--  1 ysmoon  prof       432  3  17  20:23 nchmod.c
-rwxr-xr-x  1 ysmoon  prof     6900  3  17  16:59 stat
-rw-r--r--  1 ysmoon  prof       757  3  17  16:58 stat.c
-rw-r--r--  1 ysmoon  prof         0  3  17  21:31 tempfile
-rwxr-xr-x  1 ysmoon  prof     6544  3  17  21:30 unlink
-rw-r--r--  1 ysmoon  prof       312  3  17  21:29 unlink.c
ls -l
-rwxr-xr-x  1 ysmoon  prof      6452  3  17  20:24 nchmod
-rw-r--r--  1 ysmoon  prof       432  3  17  20:23 nchmod.c
-rwxr-xr-x  1 ysmoon  prof     6900  3  17  16:59 stat
-rw-r--r--  1 ysmoon  prof       757  3  17  16:58 stat.c
-rwxr-xr-x  1 ysmoon  prof     6544  3  17  21:30 unlink
-rw-r--r--  1 ysmoon  prof       312  3  17  21:29 unlink.c
cs1 [ ysmoon {154} ~/unix/APUE ]
```

## symlink() - Symbolic Link

- Symbolic Link is an indirect pointer to a file
  - Contains a path of an actual file

```
#include <unistd.h>

int  symlink (const char *actualpath, const char *sympath );
```

- Create a Symbolic Link
  - The data of Symbolic Link File becomes actualpath
  - The size of Symbolic Link File becomes the size of actualpath string
- Returns 0 if OK, -1 on error

# File Times

- `st_atime` in `stat` structure
  - Last access time of file data (read)
- `st_mtime` in `stat` structure
  - Last-modification time of file data (write)
- `st_ctime` in `stat` structure
  - Last-change time of i-node status (`chmod`, `chown`)

# utime()

```
#include <sys/types.h>
#include <utime.h>
int utime (const char *pathname, const struct utimbuf *times );
```

- Change the access time and the modification time of a file
- If times is NULL, set as current time.
- Return 0 if OK, -1 on error
- UNIX command : touch

```
struct utimbuf {
    time_t  actime;    /* access time */
    time_t  modtime;   /* modification time */
}
```

- Each field is a value that counts seconds since 1970-1-1 00:00

# Directory

- Directory file
  - Possible to use functions such as open, read, close
  - Functions for managing a directory file are provided
- Contents of a directory file are stored as dirent type array
  - file name:
  - i-node number

```
#include <dirent.h>
struct dirent {
    ino_t d_ino;           /* i-node number */
    char  d_name[NAME_MAX + 1]; /* filename */
}
```

# Directory Access (opendir(), readdir())

```
#include <sys/types.h>
#include <dirent.h>
DIR *opendir (const char *pathname);
struct dirent *readdir(DIR *dp);
```

- opendir() : open a directory
- readdir() : read a directory contents
  - For each directory read, current file offset of the directory file is incremented by the size of dirent structure
  - For directory read, read permission of the directory is necessary.
  - Although there is write permission in directory file, write function can not be performed.  
Instead, we can use mkdir(), rmdir()
- Return the pointer address of dirent structure if OK, NULL on error

## rewinddir(), closedir()

```
#include <sys/types.h>
#include <dirent.h>

void rewinddir (DIR *dp);
int closedir (DIR *dp);
```

- `rewinddir()`
  - Move current file offset of a directory file into the start point(first entry).
- `closedir()`
  - Close directory file

# mkdir()

```
#include <sys/types.h>
#include <sys/stat.h>

int  mkdir (const char *pathname, mode_t  mode );
```

- Create a new directory
- Return 0 if OK, -1 on error
- If OK, “.” and “..” file entries are automatically created.
  - “.” points to current directory’s i-node,
  - “..” points to parent directory’s i-node



## rmkdir()

```
#include <unistd.h>

int rmkdir (const char *pathname );
```

- Remove an empty directory.
- Return 0 if OK, -1 on error

## example: listfiles.c (1/5)

```
#include <sys/types.h> // listfiles.c
#include <sys/mkdev.h>
#include <sys/stat.h>
#include <dirent.h>
#include <stdio.h>

/* typeOfFile - return the letter indicating the file type. */
char typeOfFile(mode_t mode)
{
    switch (mode & S_IFMT) {
        case S_IFREG: return('-');
        case S_IFDIR: return('d');
        case S_IFCHR: return('c');
        case S_IFBLK: return('b');
        case S_IFLNK: return('l');
        case S_IFIFO: return('p');
        case S_IFSOCK: return('s');
    }
    return('?');
}
```

## example: listfiles.c (2/5)

```
/* permOfFile - return the file permissions in an "ls"-like string. */
char* permOfFile(mode_t mode) {
    int i;
    char *p;
    static char perms[10];

    p = perms;
    strcpy(perms, "-----");
    for (i=0; i < 3; i++) {
        if (mode & (S_IREAD >> i*3)) *p = 'r';
        p++;
        if (mode & (S_IWRITE >> i*3)) *p = 'w';
        p++;
        if (mode & (S_IEXEC >> i*3)) *p = 'x';
        p++;
    }
    if ((mode & S_ISUID) != 0) perms[2] = 's';
    if ((mode & S_ISGID) != 0) perms[5] = 's';

    return(perms);
}
```

## example: listfiles.c (3/5)

```
/* outputStatInfo - print out the contents of the stat structure. */
void outputStatInfo(char *pathname, char *filename, struct stat *st) {
    int n;
    char slink[BUFSIZ+1];

    printf("%5d ", st->st_blocks);
    printf("%c%s ", typeOfFile(st->st_mode), permOfFile(st->st_mode));
    printf("%3d ", st->st_nlink);
    printf("%5d/%-5d ", st->st_uid, st->st_gid);

    if (((st->st_mode & S_IFMT) != S_IFCHR) && ((st->st_mode & S_IFMT) != S_IFBLK))
        printf("%9d ", st->st_size);
    else
        printf("%4d,%4d ", major(st->st_rdev), minor(st->st_rdev));

    printf("%.12s ", ctime(&st->st_mtime) + 4);
    printf("%s", filename);
    if ((st->st_mode & S_IFMT) == S_IFLNK) {
        if ((n = readlink(pathname, slink, sizeof(slink))) < 0)
            printf(" -> ???");
        else printf(" -> %.*s", n, slink);
    }
}
```

## example: listfiles.c (4/5)

```
int main(int argc, char **argv) {
    DIR *dp;
    char *dirname;
    struct stat st;
    struct dirent *d;
    char filename[BUFSIZ+1];
    /* For each directory on the command line... */
    while (--argc) {
        dirname = *++argv;
        if ((dp = opendir(dirname)) == NULL) /* Open the directory */
            perror(dirname);
        printf("%s:\n", dirname);
        while ((d = readdir(dp)) != NULL) { /* For each file in the directory... */
            /* Create the full file name. */
            sprintf(filename, "%s/%s", dirname, d->d_name);
            if (lstat(filename, &st) < 0) /* Find out about it. */
                perror(filename);
            outputStatInfo(filename, d->d_name, &st); /* Print out the info. */
            putchar('\n');
        }
        putchar('\n');
        closedir(dp);
    }
    return 0;
}
```

## example: listfiles.c (5/5)

### 실행 결과

```
cs1.kangwon.ac.kr - Zterm
cs1 [ ysmoon {179} ~/unix/APUE ]
cs1 [ ysmoon {179} ~/unix/APUE ] gcc -o listfiles listfiles.c
cs1 [ ysmoon {180} ~/unix/APUE ]
cs1 [ ysmoon {180} ~/unix/APUE ] ls
listfiles*      nchmod*      stat*      unlink*
listfiles.c    nchmod.c      stat.c      unlink.c
cs1 [ ysmoon {181} ~/unix/APUE ] listfiles .
..:
 2 drwxr-xr-x   2 1013/20      512 Mar 17 22:11 .
 2 drwxr-xr-x   4 1013/20      512 Mar 17 16:58 ..
 2 -rw-r--r--   1 1013/20      757 Mar 17 16:58 stat.c
14 -rwxr-xr-x   1 1013/20     6900 Mar 17 16:59 stat
 2 -rw-r--r--   1 1013/20      432 Mar 17 20:23 nchmod.c
14 -rwxr-xr-x   1 1013/20     6452 Mar 17 20:24 nchmod
 6 -rw-r--r--   1 1013/20     2547 Mar 17 22:11 listfiles.c
 2 -rw-r--r--   1 1013/20      312 Mar 17 21:29 unlink.c
14 -rwxr-xr-x   1 1013/20     6544 Mar 17 21:30 unlink
20 -rwxr-xr-x   1 1013/20     9376 Mar 17 22:12 listfiles

cs1 [ ysmoon {182} ~/unix/APUE ]
```

## chdir(), fchdir()

```
#include <unistd.h>
int chdir (const char *pathname);
int fchdir (int filedes);
```

- Change current working directory (of a process)  
(**\$cd**)
- Return 0 if OK, -1 on error
- Current working directory (cwd) is a property of a process

## getcwd()

```
#include <unistd.h>

char *getcwd (char *buf, size_t size );
```

- Get the path of current working directory (\$ **pwd**)
- Return the pointer address of buf, NULL if failed.



## example: chdir(), getcwd() (1/2)

```
#include <unistd.h> // ncd.c
#include <stdio.h>
#define PATH_MAX 1024

int main(int argc, char **argv)
{
    char path[PATH_MAX+1];

    if (argc == 1) exit(-1);

    if (argc == 2) {
        if(getcwd(path, PATH_MAX) == NULL) perror("error getcwd");
        else printf("Current working directory changed to %s \n", path);

        if(chdir(argv[1]) < 0) perror("error chdir");
        else {
            if(getcwd(path, PATH_MAX) == NULL) perror("error getcwd");
            else printf("Current working directory changed to %s \n", path);
        }
    } else
        perror("Too many arguments");
}
```

## example: chdir(), getcwd() (2/2)

```
cs1.kangwon.ac.kr - Zterm
cs1 [ ysmoon {251} ~/unix/APUE ]
cs1 [ ysmoon {251} ~/unix/APUE ]
cs1 [ ysmoon {251} ~/unix/APUE ] gcc -o ncd ncd.c
cs1 [ ysmoon {252} ~/unix/APUE ]
cs1 [ ysmoon {252} ~/unix/APUE ]
cs1 [ ysmoon {252} ~/unix/APUE ] ncd ..
Current working directory changed to /home/prof/ysmoon/unix/APUE
Current working directory changed to /home/prof/ysmoon/unix
cs1 [ ysmoon {253} ~/unix/APUE ]
cs1 [ ysmoon {253} ~/unix/APUE ]
cs1 [ ysmoon {253} ~/unix/APUE ] ncd tt
Current working directory changed to /home/prof/ysmoon/unix/APUE
Current working directory changed to /home/prof/ysmoon/unix/APUE/tt
cs1 [ ysmoon {254} ~/unix/APUE ]
cs1 [ ysmoon {254} ~/unix/APUE ]
cs1 [ ysmoon {254} ~/unix/APUE ]
```

## sync(), fsync()

```
#include <unistd.h>

void sync();
int fsync(int filedes);
```

- Write data in all the block buffers to disk
- sync() is normally called by a system daemon process for every 30 seconds
  - this guarantees regular flushing of the kernel's block buffers.
- fsync() performs sync to a single file (specified by filedes) and waits for the I/O to complete before returning.
- **O\_SYNC** flag (all write() with sync mode)